



CSC270

# Garbage Collection Fundamentals

Download class materials from  
[university.xamarin.com](https://university.xamarin.com)



**Xamarin** University

Information in this document is subject to change without notice. The example companies, organizations, products, people, and events depicted herein are fictitious. No association with any real company, organization, product, person or event is intended or should be inferred. Complying with all applicable copyright laws is the responsibility of the user.

Microsoft or Xamarin may have patents, patent applications, trademarked, copyrights, or other intellectual property rights covering subject matter in this document. Except as expressly provided in any license agreement from Microsoft or Xamarin, the furnishing of this document does not give you any license to these patents, trademarks, or other intellectual property.

© 2014-2017 Xamarin Inc., Microsoft. All rights reserved.

Xamarin, MonoTouch, MonoDroid, Xamarin.iOS, Xamarin.Android, Xamarin Studio, and Visual Studio are either registered trademarks or trademarks of Microsoft in the U.S.A. and/or other countries.

Other product and company names herein may be the trademarks of their respective owners.

# Objectives

1. Explain the Garbage Collector





# Explain the Garbage Collector

# Tasks

1. Advantages to using a GC
2. What triggers a GC?
3. Monitoring GCs in your app
4. Generational garbage collectors
5. Helping out the GC





# Memory Management

- ❖ .NET/Mono use a **Garbage Collector** (referred to as **GC**) which periodically stops your program and frees the memory your app is no longer using
- ❖ GC in Mono and .NET are similar, but the implementation is quite different, so some strategies you use in .NET may need adjustment



# Advantages to using a GC

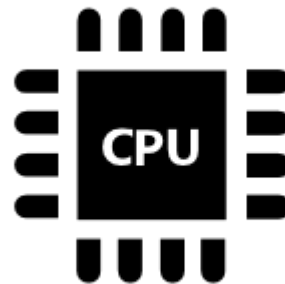
- ❖ Most modern languages/runtimes include some form of GC because of the benefits associated with it



Simple to work with



Reduces Memory Leaks



Cache Locality

# Demonstration

Show allocations and collections



# Working with the GC

- ❖ Cannot *disable* GC, but understanding how it works can help you avoid performance issues in your application, there are two things you will be concerned with



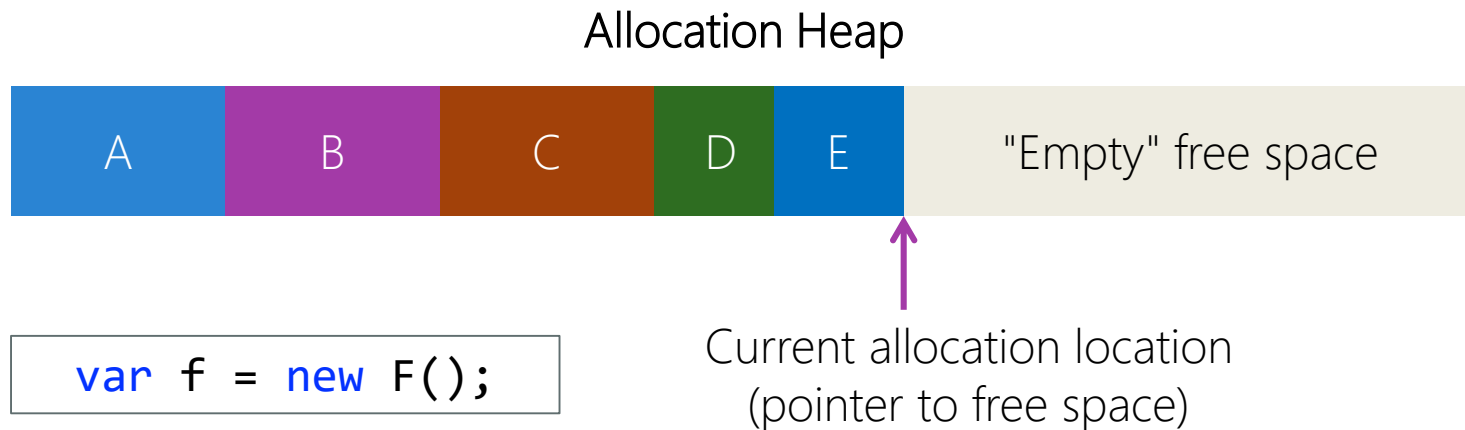
What triggers a GC?



Minimizing the pause  
time introduced by GC

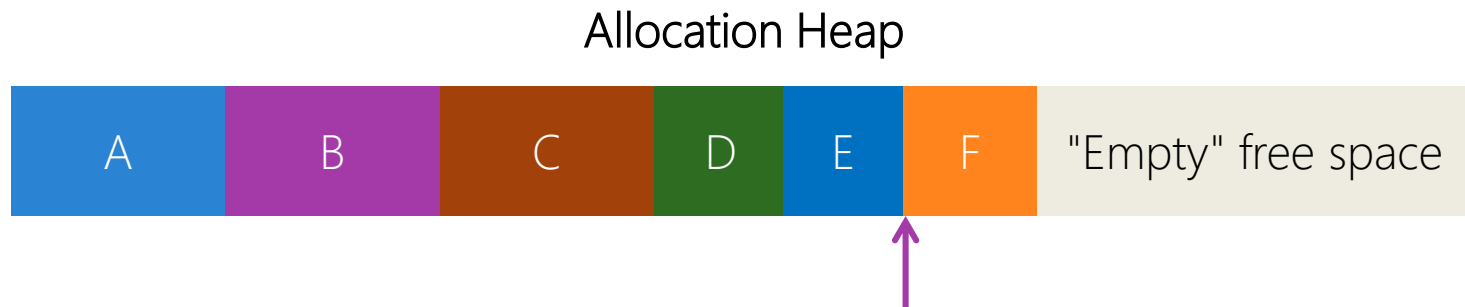
# What triggers a GC?

- ❖ GC runs as part of the memory allocation process



# What triggers a GC?

- ❖ Runtime uses a lock-free "pointer bump" allocation algorithm



Allocator returns a memory block large enough to fit an "F" instance with the block zero'd out and moves "free" pointer

New allocation location  
(pointer to free space)

# What triggers a GC?

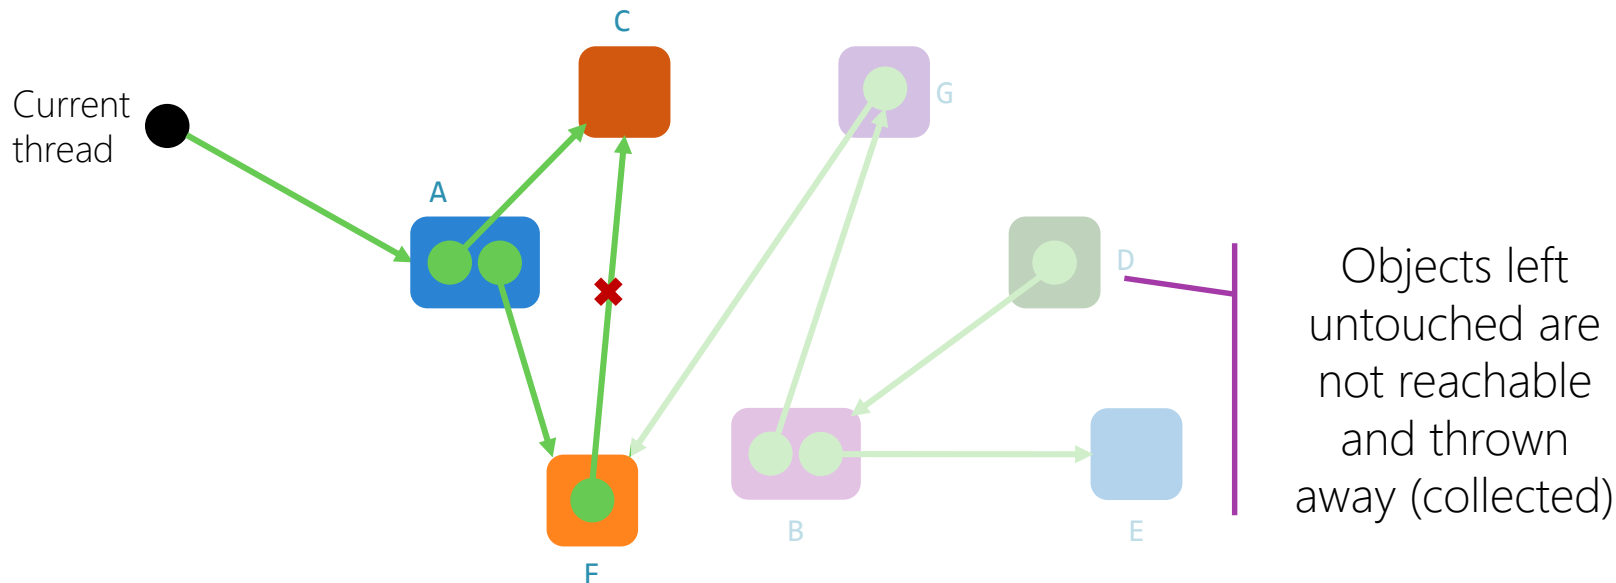
- ❖ GC kicks in on the **current thread** when there's not enough space to satisfy a memory allocation request



There is not enough memory in allocation heap for **H** object – system triggers a GC on this thread to try to free up enough space

# What happens during a GC?

- ❖ When a collection occurs, GC locates the **roots** and *traverses* the live reference graph from each one to determine what is reachable



# What is a "root"

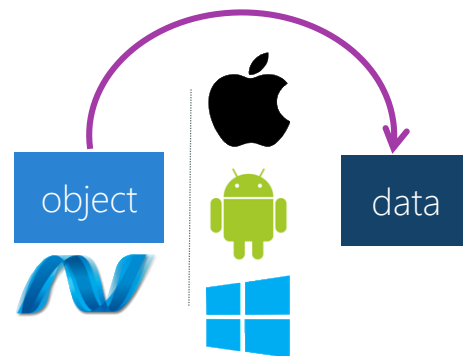
- ❖ Live references are object instances that can be accessed from a **root object** coming from one of three places:



Static property  
or field



Reference on the  
stack of a  
managed thread

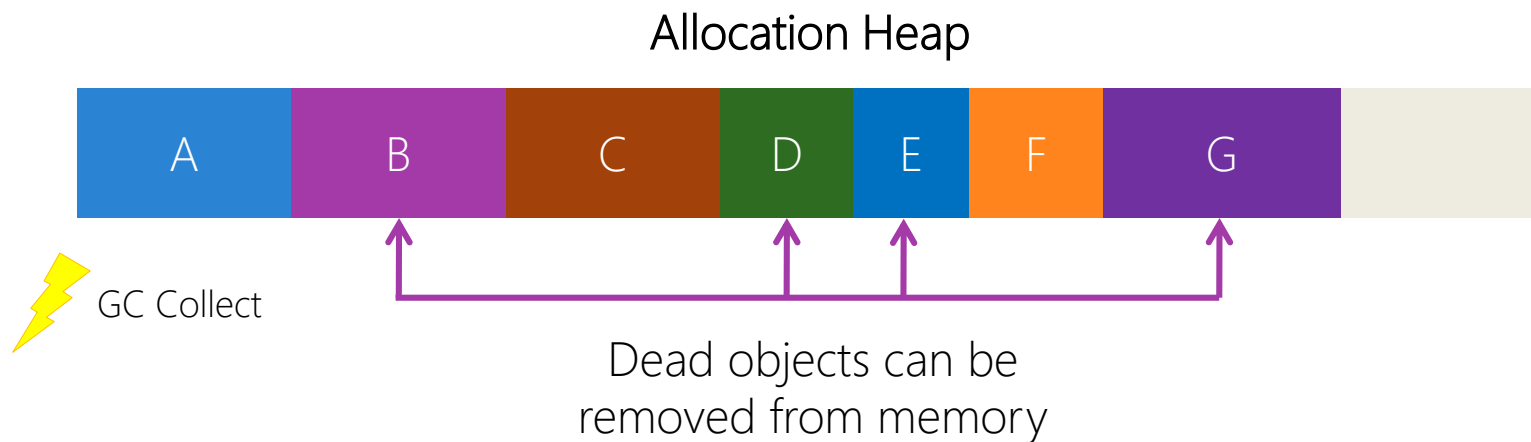


Objects passed into  
the native platform  
("pinned")



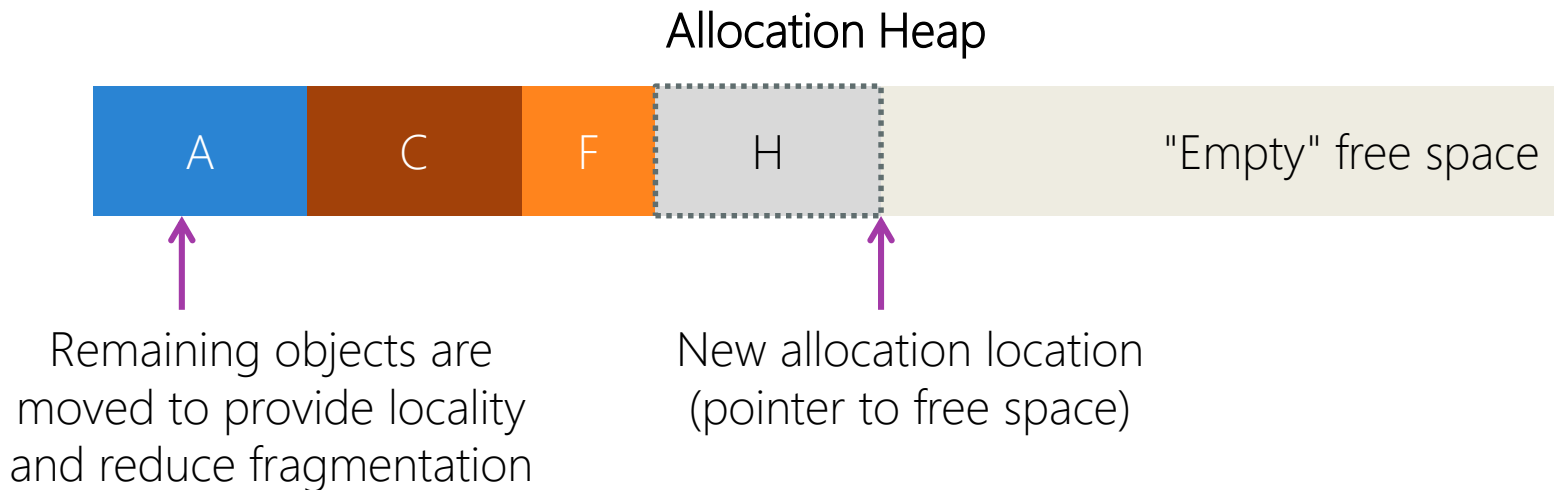
# What triggers a GC?

- ❖ GC runs and identifies **live objects**, anything not referenced can be thrown away



# What triggers a GC?

- ❖ Space is allocated for new object (H) and free space pointer is moved



# What triggers a GC?

- ❖ Application code can also *request* that a garbage collection occur

```
Image bigPicture = ...;  
...  
bigPicture = null; // Image is no longer referenced  
  
// Ask GC to run now .. to get rid of the image  
GC.Collect();
```

# Mono != .NET

❖ Mono implementation of **GC** class isn't as full featured as .NET

Does nothing

AddMemoryPressure

RemoveMemoryPressure

Partially implemented

Collect

Throws exception

WaitForFullGCApproach

WaitForFullGCComplete

RegisterForFullGC  
Notification

CancelFullGCNotification

# Flash Quiz

# Flash Quiz

- ① The runtime triggers a garbage collection \_\_\_\_\_
- a) Every two seconds on a timer
  - b) When you call **GC.Collect**
  - c) When you use the **Dispose** method on an object
  - d) When a memory allocation cannot be satisfied



# Flash Quiz

- ① The runtime triggers a garbage collection when \_\_\_\_\_
- a) Every two seconds on a timer
  - b) When you call **GC.Collect**
  - c) When you use the **Dispose** method on an object
  - d) When a memory allocation cannot be satisfied

# Flash Quiz

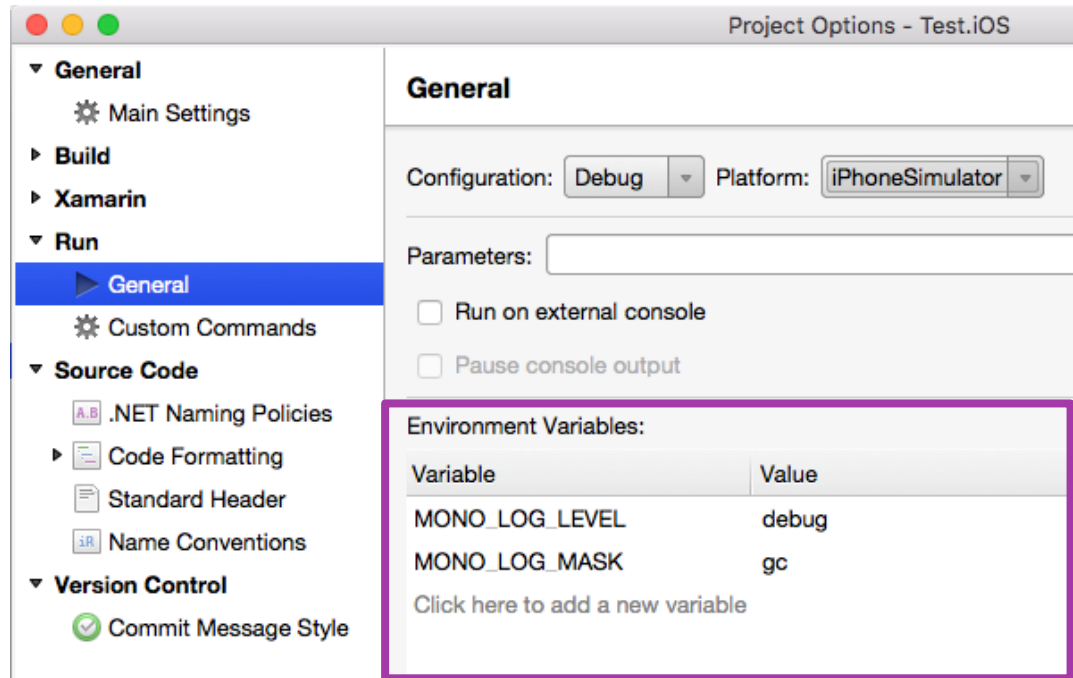
- ② GC always runs in parallel to your program code on a separate thread
- a) True
  - b) False

# Flash Quiz

- ② GC always runs in parallel to your program code on a separate thread
- a) True
  - b) False

# Monitoring GCs in your app

- ❖ Xamarin.Android outputs GC traces to the debug console for all debug builds automatically
- ❖ Xamarin.iOS requires you add two environment variables to get the output in the debug console



# Debug output

- ❖ GC messages are pre-pended with **GC\_** tags

What triggered this GC



GC\_MINOR: (Nursery full) pause 0.42ms, total 0.43ms, bridge 0.00ms promoted 5K major 5728K los 37K

GC\_MAJOR: (user request) pause 6.07ms, total 6.07ms, bridge 0.00ms major 5728K/5728K los 29K/29K



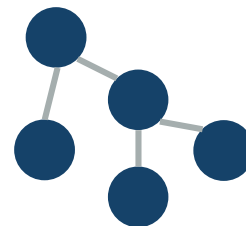
How long did the GC take

# How long does a GC take?

- ❖ The duration (pause time) it takes to do a collection depends on two primary factors:



What type, or **generation** of collection is occurring

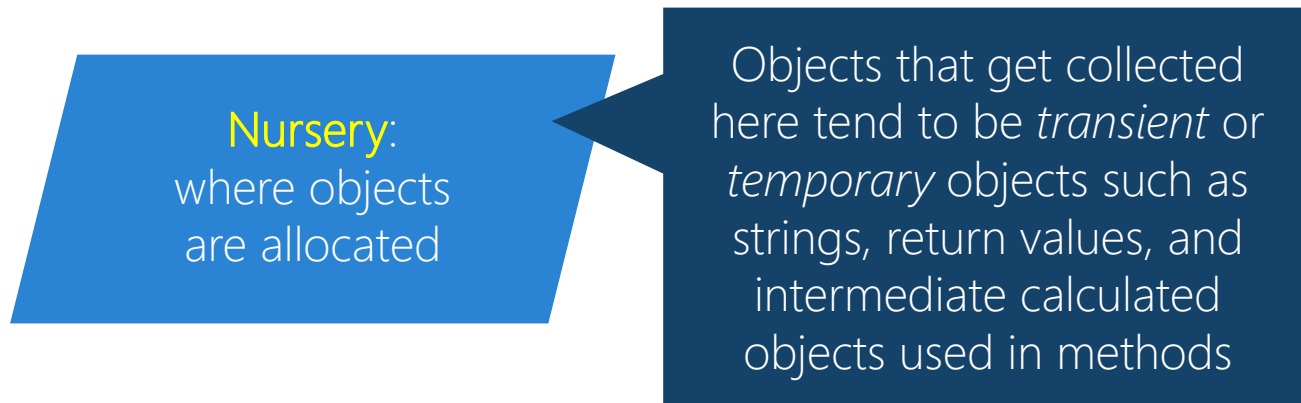


How many objects are reachable and have to be examined



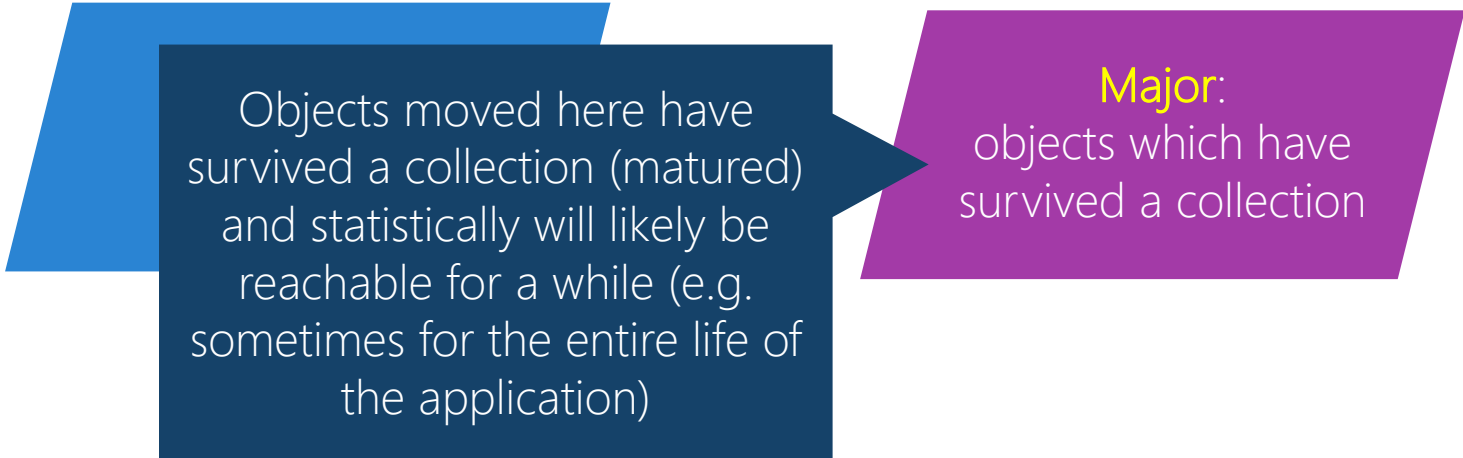
# Generational garbage collectors

- ❖ GC works off the premise that *most objects die young*; Mono and Android separate objects into two *generations* (.NET uses three)



# Generational garbage collectors

- ❖ GC works off the premise that *most objects die young*; Mono and Android separate objects into two *generations* (.NET uses three)

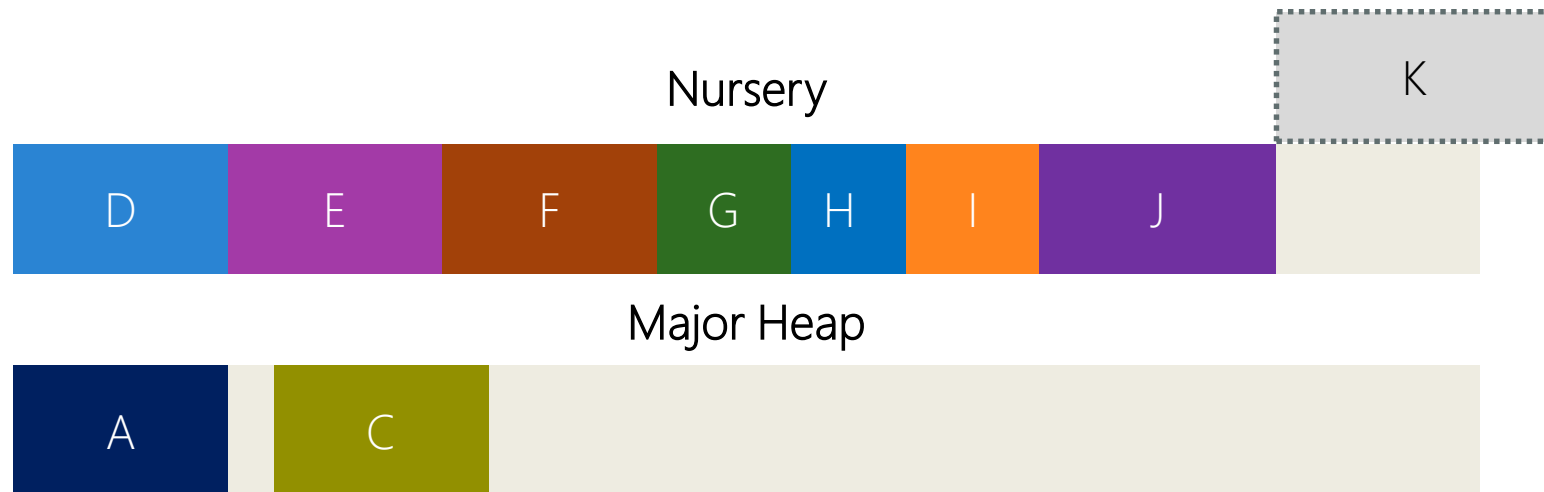
A diagram illustrating the two generations of garbage collection. It consists of two overlapping parallelogram shapes. The left shape is dark blue and contains text describing objects that have survived a collection and are likely to be reachable for a while. The right shape is purple and contains text describing major objects that have survived a collection. A blue arrow points from the dark blue shape to the purple shape, indicating a transition or relationship between the two generations.

Objects moved here have survived a collection (matured) and statistically will likely be reachable for a while (e.g. sometimes for the entire life of the application)

**Major:**  
objects which have survived a collection

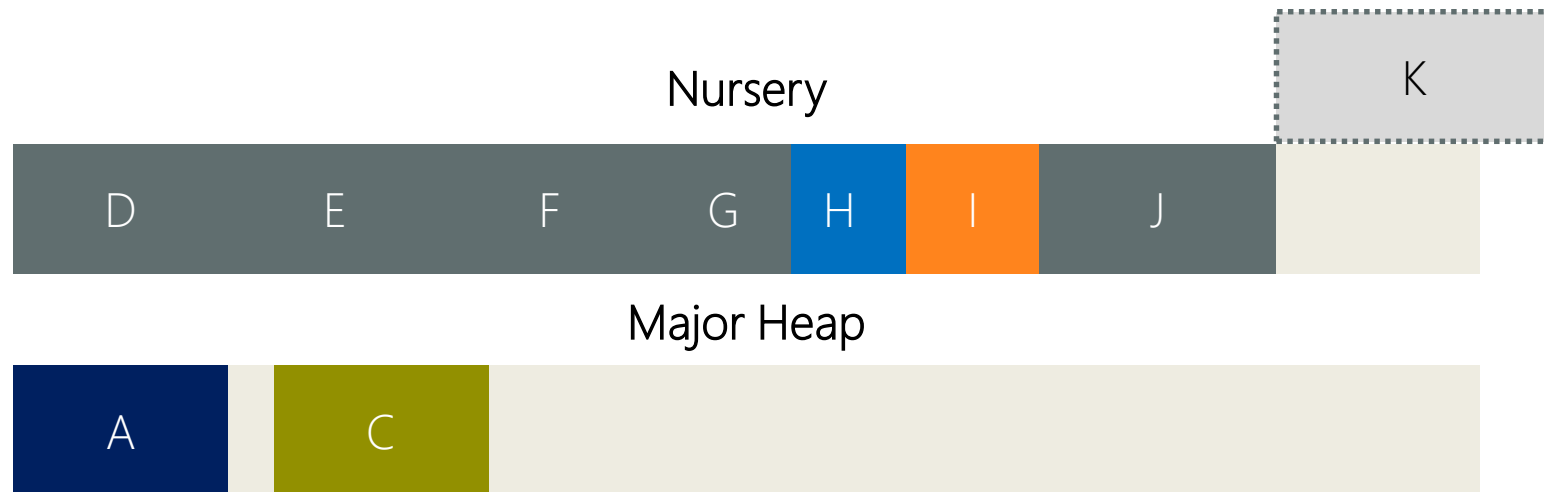
# Collecting the nursery

- ❖ Nursery collections happen anytime an allocation fails; since the nursery is small, this happens frequently and pauses tend to be short



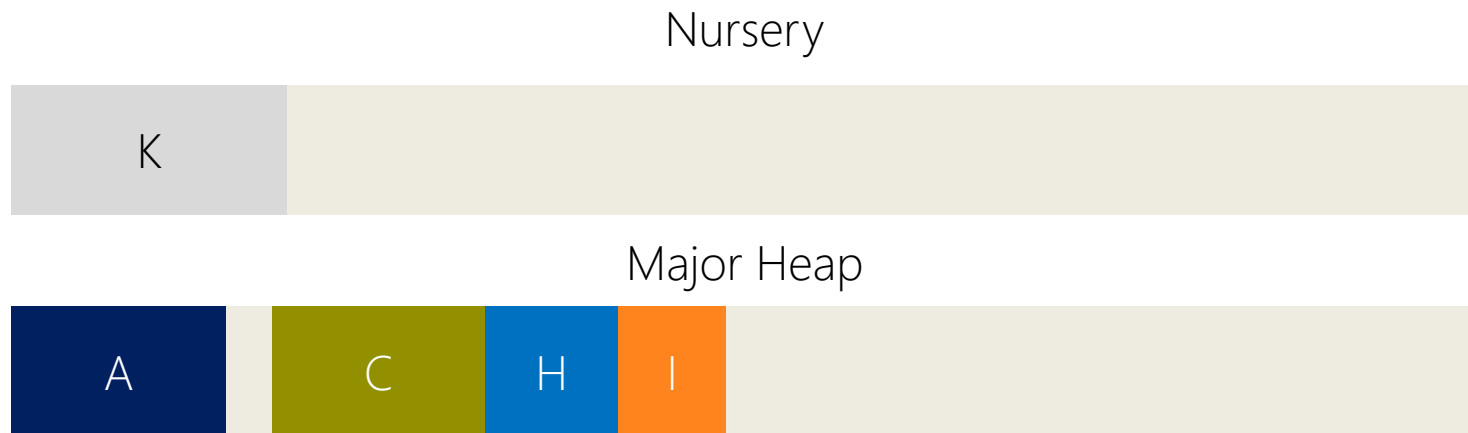
# Collecting the nursery

- ❖ Nursery collections happen anytime an allocation fails; since the nursery is small, this happens frequently and pauses tend to be short



# Adding generations

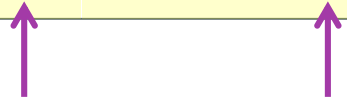
- ❖ Live objects are moved into **major heap** and pointer references are changed to point to the new locations



# Debug log: promoting memory

- ❖ **GC\_MINOR** reports promotion size and current major generation size

GC\_MINOR: (Nursery full) pause 0.42ms, total 0.43ms, bridge 0.00ms promoted 5K major  
5728K los 37K



How much memory  
survived nursery collection  
and was copied

How much is currently in  
the major generation



# Debug log: memory growth

- ❖ **GC\_MAJOR** reports current and previous memory usage for major heap and LOS; watch these values to identify potential leaks

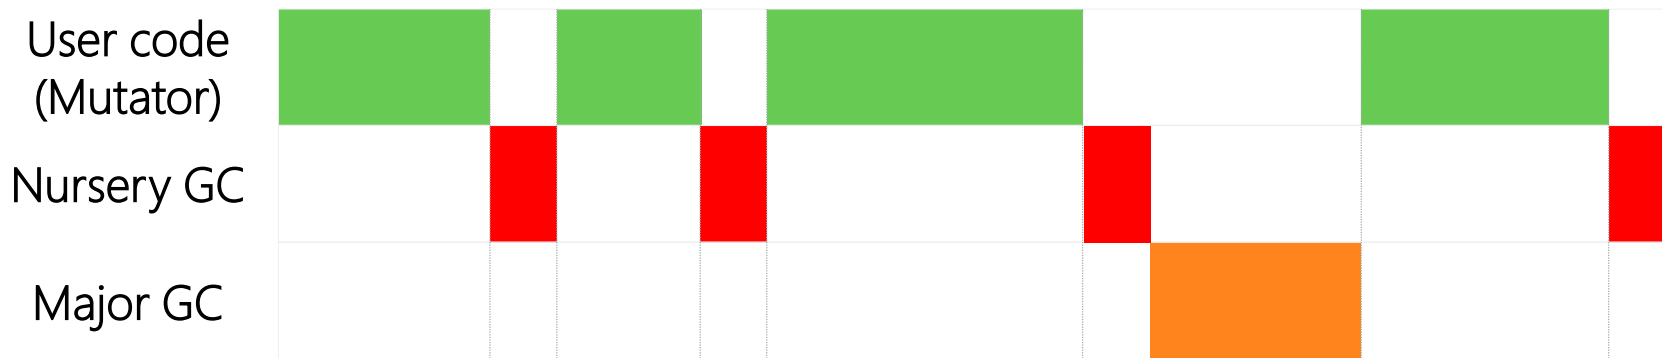
GC\_MAJOR: (user request) pause 6.53ms, total 6.54ms, bridge 0.00ms major  
1008K/16912K los 64K/0K

How much memory is  
currently in use by the  
major generation

How much memory *was* in  
use prior to this collect?

# Collecting the major heap

- ❖ Major heap is collected less frequently; typically due to promotion growth or when there's not enough space in the heap to complete a nursery collection



**Remember:** each time the GC runs, it **stops the world** and our app is not executing user code, major collections could potentially stop your app for up to a second

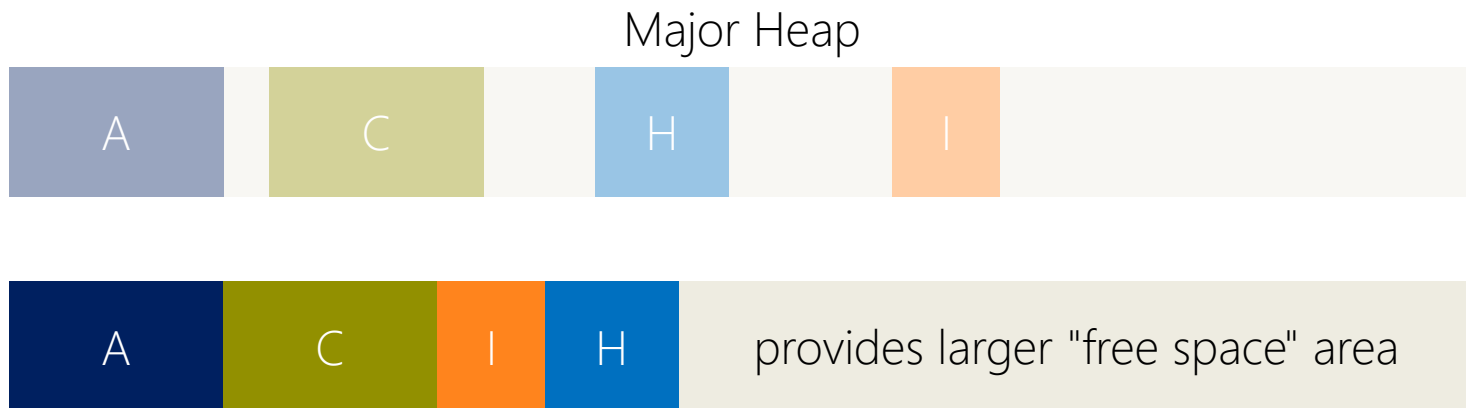
# Nursery vs. Major collection ratio

- ❖ System allows the major heap to **grow by ~1/3** before initiating a full collection
- ❖ Healthy GC behavior is around **1:10** for major to minor collections, depending on what the app does
- ❖ Check debug log or use **GC.CollectionCount** to determine the ratio for your app

```
45.571 GC_MAJOR: ...
45.588 GC_MINOR: major 4256K
45.606 GC_MINOR: major 6320K
45.625 GC_MINOR: major 7888K
45.641 GC_MINOR: major 8336K
45.660 GC_MINOR: major 10416K
45.682 GC_MINOR: major 11840K
45.699 GC_MINOR: major 12160K
45.719 GC_MINOR: major 14240K
45.738 GC_MINOR: major 15536K
45.754 GC_MINOR: major 15728K
45.772 GC_MINOR: major 17808K
45.790 GC_MINOR: major 18976K
45.806 GC_MINOR: major 19040K
45.829 GC_MINOR: major 21104K
45.851 GC_MAJOR: ...
```

# Compacting the major heap

- ❖ Major heap is organized in *buckets* and objects are moved from the nursery into a specific bucket based on size; GC will *periodically compact* these buckets to reduce fragmentation

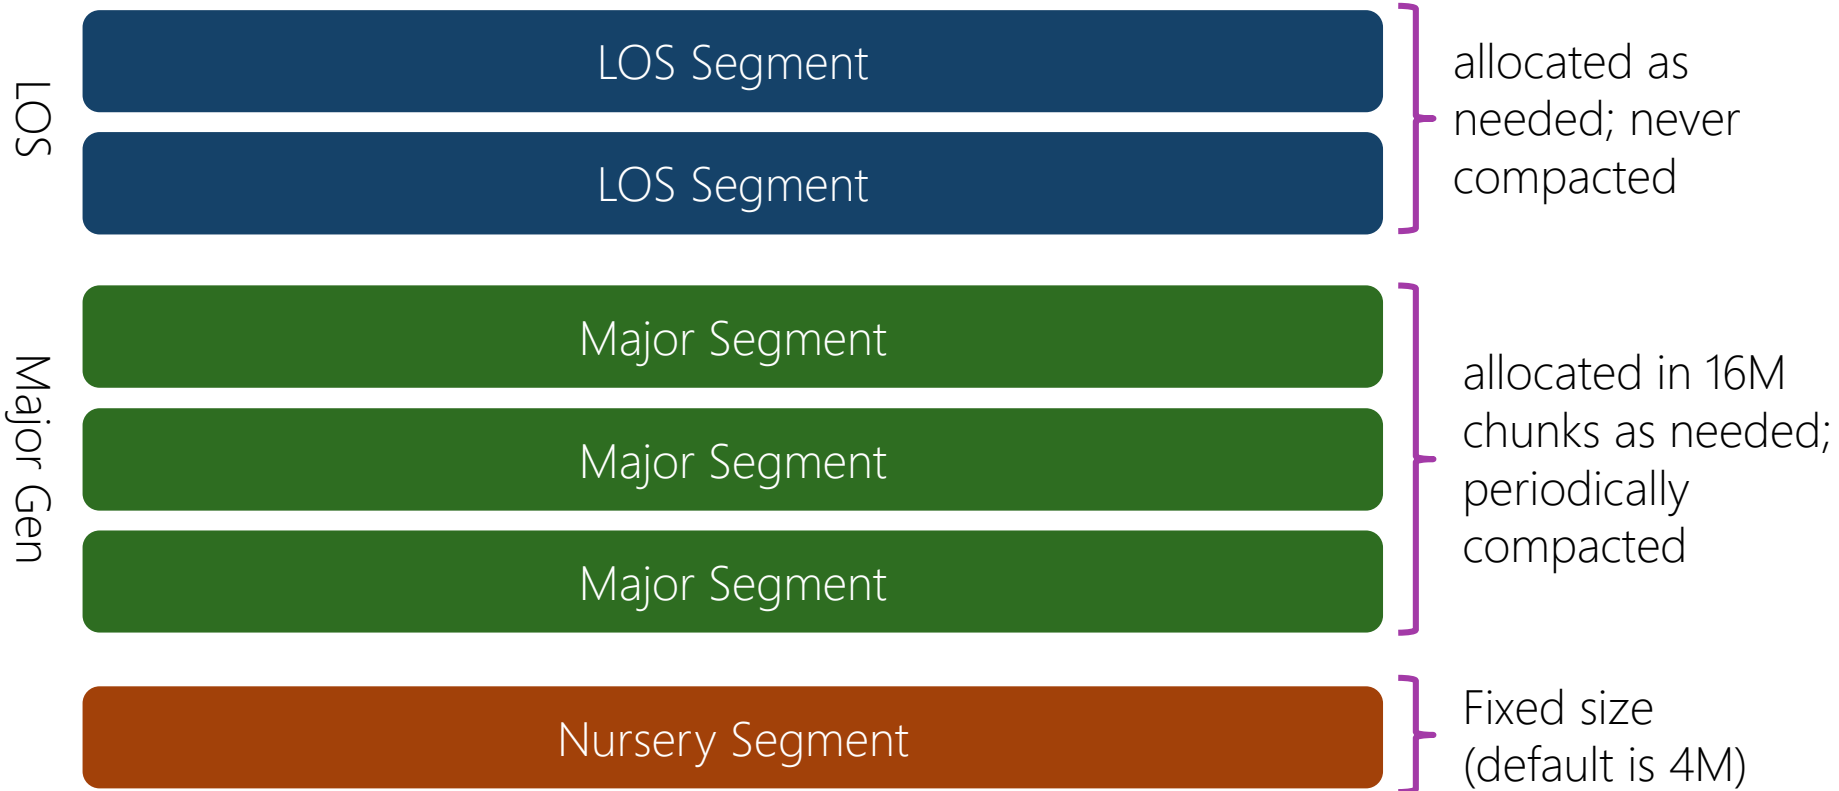


# Big objects are different

- ❖ Some objects are more expensive to move/compact due to their size (i.e. large arrays)
- ❖ These *large objects* are placed into a special heap referred to as the **Large Object Space (LOS)** (or Large Object Heap in .NET)
- ❖ Collected on full collections and not compacted



# SGen memory management



# Flash Quiz

# Flash Quiz

- ① The **nursery generation** will grow and shrink based on usage
- a) True
  - b) False



# Flash Quiz

- ① The **nursery generation** will grow and shrink based on usage
- a) True
  - b) False

# Flash Quiz

- ② To turn on GC diagnostic tracing in iOS you need to \_\_\_\_\_
- a) Do nothing. It's on by default.
  - b) Add an entry to your **info.plist**
  - c) Add two environment variables to the project settings

# Flash Quiz

- ② To turn on GC diagnostic tracing in iOS you need to \_\_\_\_\_
- a) Do nothing. It's on by default.
  - b) Add an entry to your `info.plist`
  - c) Add two environment variables to the project settings

# Setting variables to null

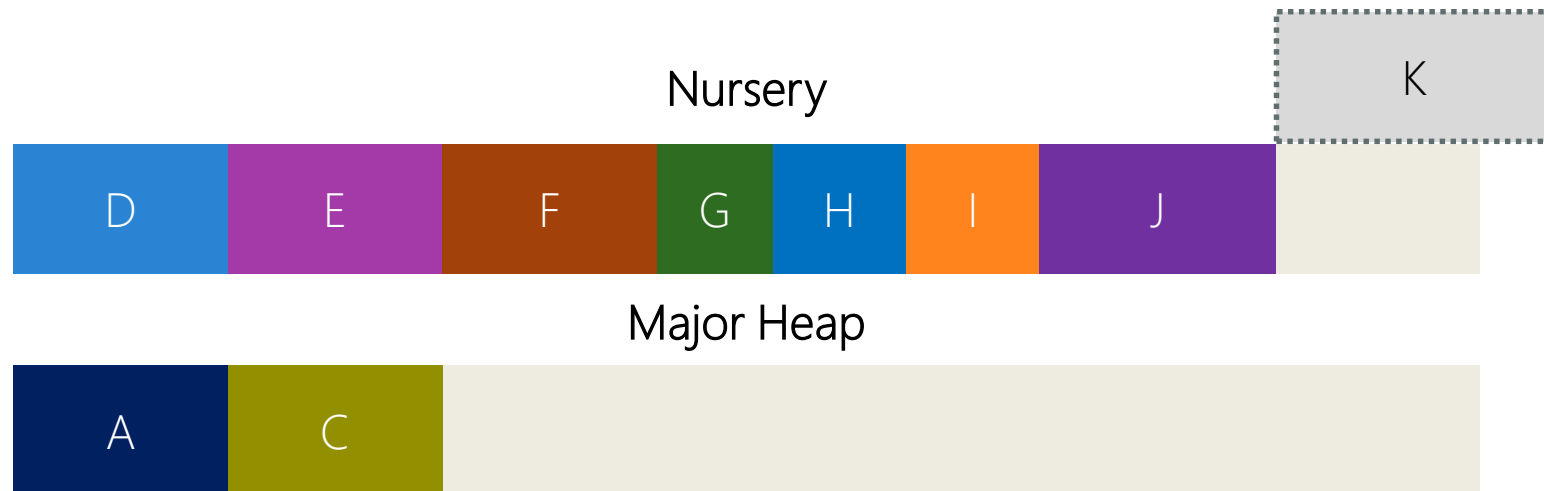
- ❖ Can set **static** or class instance fields and properties to **null** to let the GC collect the associated references

```
public class App : Application
{
    public IDBFactory DbFactory { get; set; }

    protected override void OnStart () {
        IDBRepository repo = DbFactory.Create ();
        ...
        // Don't need factory anymore.
        DbFactory = null;
    }
}
```

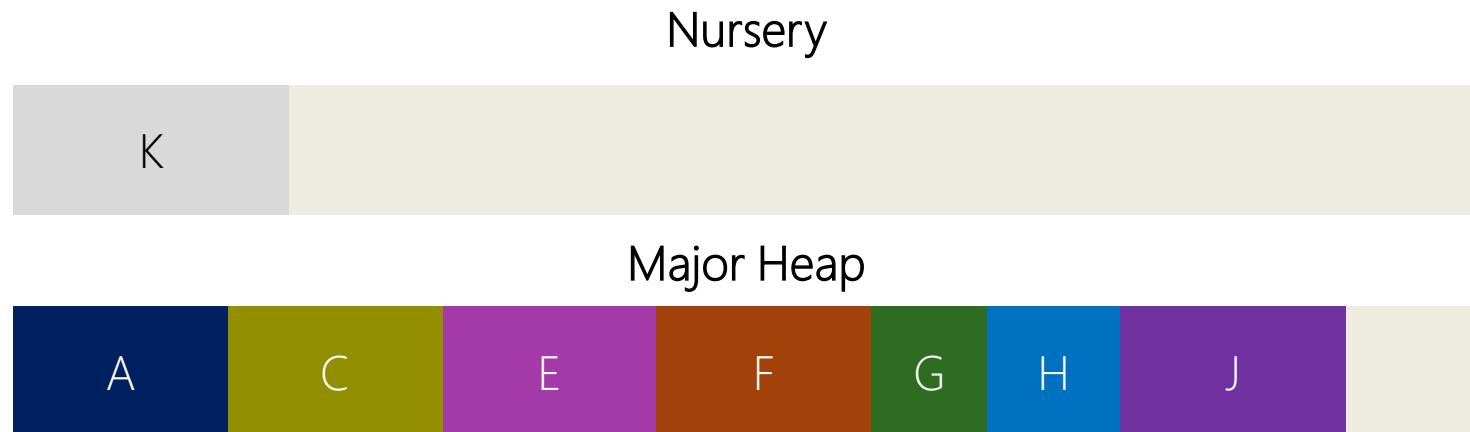
# Best case scenario

- ❖ We want our objects to *either* have very short lifetimes and be removed in the nursery collection, or to live long lifetimes



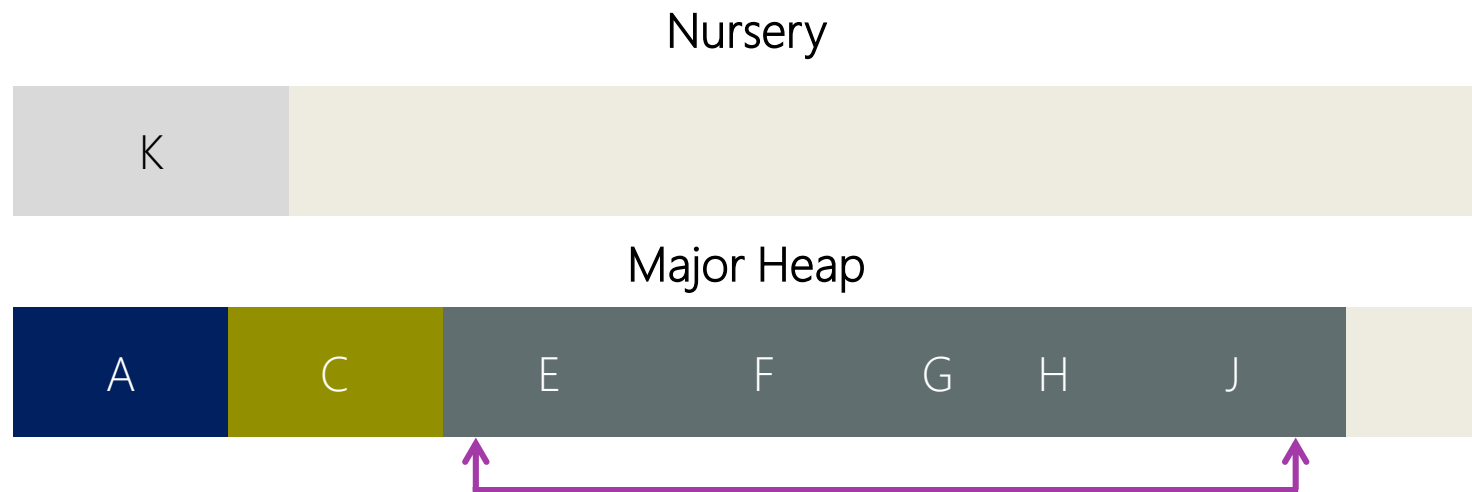
# Worst case scenarios

- ❖ App is executing code which still had references to a bunch of temporary objects when a nursery collection occurs, which moves *all of these objects to the major generation!*



# Worst case scenarios

- ❖ The worst thing scenario is when object(s) are promoted to the major generation *and then expires shortly thereafter*



Now these objects require a major generation collection to go away ...

# Minimizing the mid-life collection

- ❖ Can call **GC.Collect** before starting memory intensive data processing to try to avoid mid-processing collections

```
void DoSomeExpensiveThing()
{
    // Force a collect before we start so nursery is empty
    GC.Collect ();
    ... // Do work that allocates objects
}
```



# Tweaking the GC

- ❖ **MONO\_GC\_PARAMS** environment variable can be set to tweak the GC settings based on your app needs



nursery-size=#bytes  
(defaults to 4M)

Can change the fixed segment size used for the nursery to minimize the # of collections at the expensive of a larger working set

```
MONO_GC_PARAMS=nursery-size=8M
```

# Tweaking the GC

- ❖ **MONO\_GC\_PARAMS** environment variable can be set to tweak the GC settings based on your app needs



nursery-size=#bytes  
(defaults to 4M)



soft-heap-limit=#bytes  
(defaults to -1)

Can set a max working set limit to reduce how often the major heap is collected

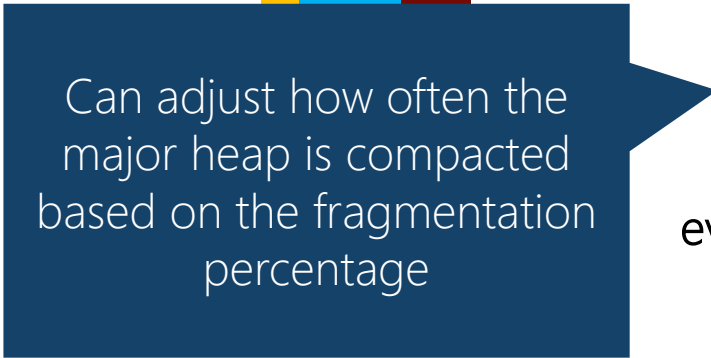
```
MONO_GC_PARAMS=soft-heap-limit=512M
```

# Tweaking the GC

- ❖ **MONO\_GC\_PARAMS** environment variable can be set to tweak the GC settings based on your app needs



nursery-size=#bytes  
(defaults to 4M)



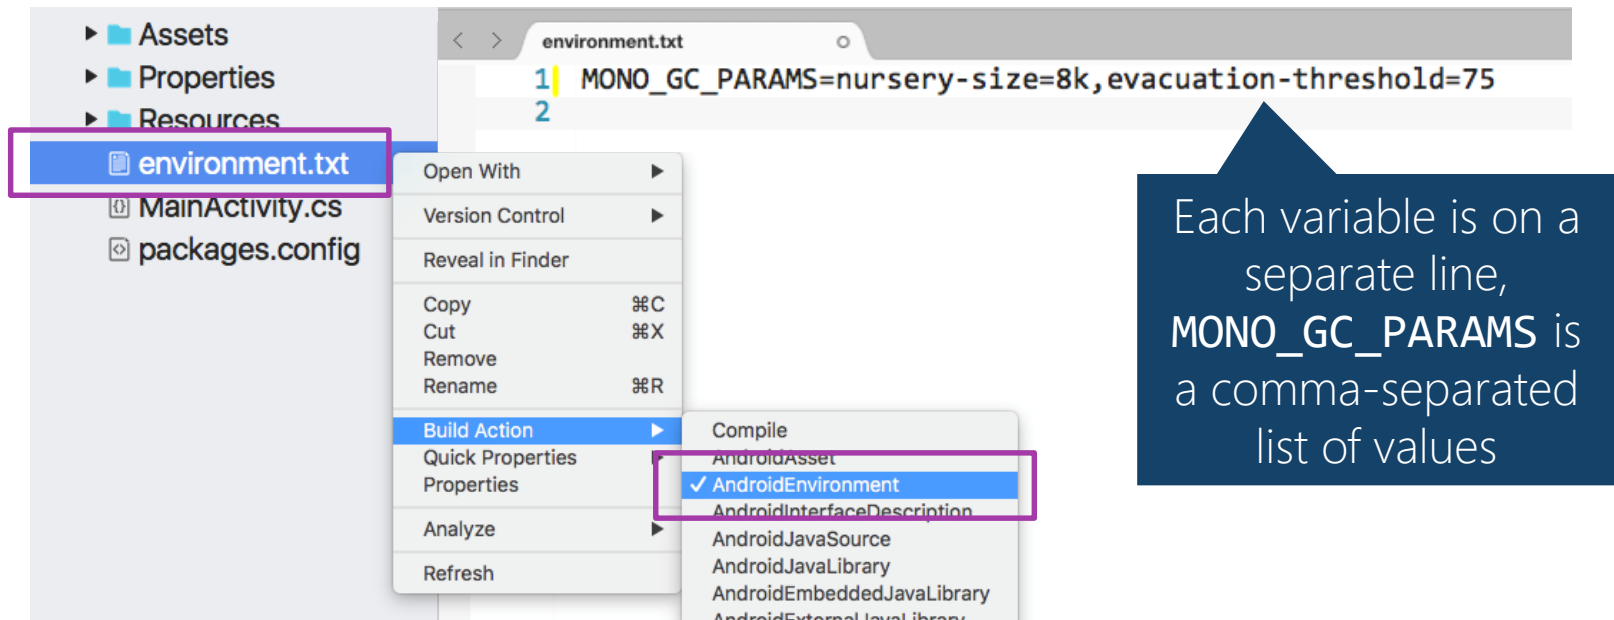
Can adjust how often the major heap is compacted based on the fragmentation percentage



evacuation-threshold=%  
(defaults to 66)

```
MONO_GC_PARAMS=evacuation-threshold=50
```

# Setting environment variables in Android



The screenshot shows an IDE interface. On the left, a file explorer lists 'Assets', 'Properties', 'Resources', 'environment.txt' (highlighted with a purple box), 'MainActivity.cs', and 'packages.config'. The main editor displays 'environment.txt' with the following content:

```
1 MONO_GC_PARAMS=nursery-size=8k,evacuation-threshold=75
2
```

A context menu is open over the file, with 'Build Action' selected. A sub-menu is also open, showing 'AndroidEnvironment' selected (highlighted with a purple box). Other options in the sub-menu include 'Compile', 'AndroidAsset', 'AndroidInterfaceDescription', 'AndroidJavaSource', 'AndroidJavaLibrary', 'AndroidEmbeddedJavaLibrary', and 'AndroidExternalJavaLibrary'.

Each variable is on a separate line,  
**MONO\_GC\_PARAMS** is  
 a comma-separated  
 list of values

# Demonstration

Adjust the nursery size



**Xamarin**  
University

# Summary

1. Advantages to using a GC
2. What triggers a GC?
3. Monitoring GCs in your app
4. Generational garbage collectors
5. Helping out the GC



# Thank You!

Please complete the class survey in your profile:  
[university.xamarin.com/profile](https://university.xamarin.com/profile)

